
Distance to the Boundary Is Not Steering Resistance

Alignment widens the behavioural margin, not the steering lever, so fixed-dose steering audits misrank aligned models

Anonymous submission

Abstract

A refusal direction fit on a base language model is widely reported to “go stale” on the model’s aligned descendants: ablate the direction and the base model complies, but the aligned model still refuses. We show the direction does not weaken. The target moves. Across the full OLMo-3 training flow the model’s behavioural margin, its unsteered distance from the refusal decision boundary, grows $9.5\times$, from 0.81 to 7.75 logits, while the direction’s per-dose steering efficacy stays nearly constant (coefficient of variation 0.13, a $1.5\times$ spread whose maximum is the aligned model itself; a gradient-trained direction, six times more efficient and near-orthogonal, is similarly near-constant, CV 0.196). A fixed-magnitude intervention therefore fails on aligned models for one banal reason: the same push is now small against a far wider margin. This inverts a natural ranking. Under fixed ablation and fixed-dose steering the aligned model looks least controllable, crossing 1 % of prompts where the base model crosses 92 %, yet its per-dose lever is the strongest in the set. Honesty supplies a convergent control: its margin barely grows ($1.2\times$), and its audit artifact nearly vanishes, so margin growth drives the misranking rather than any property of the lever. The mechanism replicates across four model families (OLMo-3, Qwen3-8B, Llama-3.1-8B, Gemma-2-9B), where the margin grows between $3.7\times$ and $9.8\times$. Margin growth alone is enough to cause the misranking, whatever the direction’s own steering power does: it is non-weakening on OLMo, roughly constant on Gemma, and grows on Qwen and Llama. The effect holds at larger scale (OLMo-3 32B). A second, independent warning follows: onset-control is not harm-control. The base direction still flips an aligned model’s refusal onset to compliance, but the peak HarmBench-judged genuine harm achievable across the dose sweep decays with alignment (0.80 to 0.05 on OLMo, replicated on Qwen; this decay is confounded with dose-driven degeneration and is a hypothesis pending a coherence-matched measurement, §8), and prefix-based refusal metrics increasingly overstate compliance, their correlation with true harm flipping from $+0.96$ to -0.53 . Safety audits that ablate a fixed direction understate aligned-model controllability; audits that read refusal-onset as harm overstate it.

1 Introduction

Activation steering finds a residual-stream direction that correlates with a behaviour and adds or subtracts it, and it is a standard tool for interpreting and controlling language models. A consequential question follows for anyone auditing a deployed model: does a direction found on an earlier checkpoint still control the behaviour after the model is aligned? If white-box safety interventions built on base-model analysis silently expire during alignment, they expire when they matter.

The published answers disagree. One line of work reports that base-derived steering vectors remain effective on aligned instruct models [7]. Another reports that post-training directions steer better than pretraining ones [8]. A third, from our own earlier experiments with the single-direction ablation

method of Arditi et al. [1], finds that the base refusal direction loses its grip: ablate it and the aligned model keeps refusing. Three groups, three answers, one model family.

We resolve the disagreement with one measurement and one distinction. Every prior study fixes the *magnitude* of the intervention and reads the *outcome*. But a fixed intervention is not a fair probe of two models whose behaviour sits at different distances from the decision boundary. An aligned model refuses hard; its margin is large; a fixed nudge moves it far and still leaves it refusing. We separate the **lever** (displacement produced per unit dose) from the **load** (the margin the lever must overcome), measure both across the OLMo-3 flow, and find the lever essentially constant while the load grows an order of magnitude. Honesty, a concept whose load does not grow, supplies a convergent control. A validity check with an independent harm judge then reveals a second effect the onset-level view hides: the direction keeps flipping the aligned model’s onset long after it stops eliciting genuine harm.

Contributions.

- We show that fixed-magnitude steering and ablation audits systematically misrank aligned models as less controllable, and identify the cause: alignment widens the behavioural margin ($3.7\times$ – $9.8\times$ across four model families) while the direction’s per-dose lever does not weaken — a non-inferiority test confirms this at all nine post-training checkpoints (§3–§4).
- We establish this with a control concept: honesty, whose margin barely grows ($1.2\times$), shows essentially no audit artifact; margin growth, not the lever, drives the misranking (§5).
- We report a second, opposite audit failure: onset-control is not harm-control. The base direction still flips an aligned model’s refusal onset, but genuine (HarmBench-judged) harm decays with alignment, and prefix-based refusal metrics increasingly overstate compliance (§6).
- We give the honest scope: the per-dose lever is non-weakening on OLMo (non-inferiority holds) but not general off it (approximate on Gemma, growing on Qwen and Llama); the universal, sufficient cause of the misranking is margin growth. All code, pinned commits, split fingerprints, and a discarded-statistic autopsy are released.

2 Setup

Models. The OLMo-3 7B family [11], whose entire training flow is public: base; Think lineage at SFT 1k/15k/43k, DPO, first and last RLVR steps; Instruct; and two RL-Zero variants (RL applied directly to the base, no SFT). For generalization: OLMo-3 32B (base/SFT/DPO/RLVR-last) and three further base→instruct families, Qwen3-8B, Llama-3.1-8B, and Gemma-2-9B (the latter two via full-precision community mirrors, verified bf16 with correct parameter counts, as the official repos are license-gated). Every checkpoint is pinned by commit hash.

Directions. A refusal direction at layer 20 (proportional layer for 32B/Qwen), fit once on the base model as the difference in mean residual-stream activation between 200 harmful [12] and 200 benign [10] prompts under a neutral scaffold applied identically to every checkpoint, and carried unchanged to every descendant (this is the staleness question). We fit it three ways (difference-in-means, logistic regression at $\text{tol} = 10^{-10}$, and a gradient-trained steering vector), which disagree geometrically (cosine 0.16–0.74) yet agree on every conclusion. The honesty direction is fit the same way, on Azaria–Mitchell true/false statements [2].

Lever and load. Steering uses $h_l \leftarrow h_l + c \mu_l \hat{v}$ (input-normalised, comparable across checkpoints; the dose scale μ_l is itself nearly constant across the flow, CV 0.035, so efficacy per unit c and per unit activation-norm coincide). The readout is the concept’s logit gap (refusal- vs compliance-onset tokens; ids fixed). The *load* (margin) is the mean unsteered gap. The *lever* (efficacy) is the slope of mean per-prompt displacement against $|c|$ in the near-zero regime (displacement saturates at large $|c|$; we use the linear region). We also measure behavioural efficacy (refusal-rate change vs dose on real generations) and validate genuine harm with the HarmBench classifier [6].

Rigour. Every headline number carries a 95 % bootstrap CI (1000 resamples; paired for comparisons) and a 20-direction random-direction null. Splits are fingerprinted; every refusal run reproduces

Table 1: Three instruments, three rankings of the same four models. Fixed-magnitude instruments rank the aligned model last; per-dose efficacy ranks it first. The margin column is the load, shown for reference, not an instrument.

model	fixed ablation	fixed dose	efficacy (lever)	margin
base	1.00	0.92	3.84	0.81
Instruct	0.33	0.01	5.61	7.75
RL-Zero Math	1.00	0.88	4.01	0.99
RL-Zero Code	1.00	0.93	4.12	0.83




Figure 1: Three instruments, three rankings of the same four models (a single slopegraph; each vertical position is min–max normalised within its instrument). Left and middle—fixed-magnitude ablation and fixed-dose steering—place the aligned model *last*, the basis for “base directions go stale”; right—per-dose efficacy—places it *first*. The disagreement is instrumental: the first two confuse distance-to-boundary (load) with resistance-to-being-moved (lever).

the same fingerprint, and the base model reproduces bit-for-bit across sessions. Five pipeline smoke tests pass. We report where a claim is clean (OLMo) and where only approximate.

3 The instrument decides the answer

Rank the same four models (base, Instruct, two RL-Zero variants) by “how steerable is refusal,” and the ranking depends entirely on the instrument. Reading each instrument’s native metric off the same sweep (Table 1), fixed-magnitude ablation and fixed-dose steering, the two instruments the prior literature used, rank the aligned Instruct model last: it crosses 1 % of prompts where base crosses 92 %. Yet its per-dose lever is the strongest of the four. The instrument that reads “uncontrollable” is measuring the model with the most intact lever. “Controllable” here is onset-level: §6 shows that flipping the aligned model’s onset does not by itself elicit genuine harm, so this concerns the direction’s grip on the first-token disposition, not how dangerous the aligned model can be made.

4 The lever holds; the load grows

Measuring lever and load at ten points across the flow (Figure 2), the displacement produced per unit dose ranges only 3.6–5.6 (a $1.5\times$ spread, CV 0.13) across the ten checkpoints (base, six Think post-training stages, Instruct, both RL-Zero variants), clearing the random-direction null everywhere (≈ 4 –7). The residual dispersion is systematic, not noise: the aligned model has the *largest* lever, so we call the lever non-weakening and mildly increasing for refusal, and reserve “invariant” for the honesty control (CV 0.021). The per-layer L20 CV of 0.19 reported below is a distinct quantity,




Figure 2: The lever changes little; the load grows an order of magnitude. Across the staged flow base $\rightarrow$ SFT $\rightarrow$ DPO $\rightarrow$ RLVR $\rightarrow$ Instruct, the behavioural margin (load) grows $9.5\times$ (to 7.75 logits) while steering efficacy (lever) rises only to $1.46\times$ of its base value. Over all ten OLMo checkpoints (the eight shown plus the two RL-Zero variants) efficacy has CV 0.13. A fixed-magnitude ablation delivers a fixed displacement, so it fails on aligned models chiefly because the margin widened. The lever line shows 95% CIs; the margin fold is $9.5\times$ (95% CI [7.4, 13.7]).

over a three-checkpoint layer-robustness subset, not this ten-point flow. A non-inferiority test with a pre-stated equivalence margin ($\Delta = 0.20\times$ base efficacy = 0.77; paired bootstrap over prompts, 1000 resamples) establishes the lever as non-weakening at all nine post-training checkpoints: the worst case (SFT 43k) has paired difference vs base -0.21 (95% CI $[-0.30, -0.12]$), inside the margin, and the aligned Instruct model’s lever is significantly *above* base ($+1.77$, 95% CI $[+1.64, +1.91]$), which is why we say mildly increasing rather than flat. The ten-checkpoint efficacy CV is 0.127 (95% CI $[0.119, 0.137]$). Per-checkpoint efficacy CIs are in Table 4 (Appendix A) and drawn on Figure 2. The margin fold is $9.5\times$ (95% CI [7.4, 13.7]).

The near-zero slope is a local quantity, but the fixed-dose audit acts at large dose where displacement saturates. We therefore also measure the saturation ceiling (maximum achievable mean displacement) and the secant lever to each model’s own crossing dose. The ceiling exceeds the margin at every checkpoint (Instruct’s is the largest, 11.2 vs its margin 7.75), so every model can in fact be crossed: the aligned model does, behaviourally (refusal 1.00 to 0.18 at its own dose, §6). The fixed-dose audit fails because a fixed small push lands short of a wider margin, not because saturation blocks it, and the secant lever to the crossing dose is comparable at the endpoint (Instruct 5.99 vs base 6.42). Over the same span the margin climbs $9.5\times$, from 0.81 to 7.75 logits. A fixed intervention delivers a fixed displacement (efficacy $\times$ dose); as the load grows tenfold while that displacement stays fixed, it crosses fewer prompts. The direction is not going stale. The target is receding.

This near-constancy is not an artifact of how we fit the direction. A gradient-trained steering direction, six times more efficient (mean 26.1 vs 4.48 over the four E1 checkpoints) and near-orthogonal to diff-in-means (cosine 0.21), is similarly near-constant (efficacy CV 0.196, against 0.13 for difference-in-means). The near-constancy holds across mid-to-late layers (per-layer CV: L20 0.19, L24 0.11, L28 0.07; early layers reorganise with alignment), and at 32B scale (CV 0.041, margin $2.3\times$). It is a property of the direction as a control axis, not of the estimator, the layer, or the model size.

Reconciliation. “Base vectors remain effective”: the lever is near-constant. “They go stale under fixed ablation”: the load grew, so a fixed removal no longer reaches the boundary. “Post-training directions steer better”: refit at a later checkpoint, a direction is simply calibrated to that checkpoint’s larger load. Three reports, one near-constant lever, seen through instruments that read the load as the lever.

Table 2: The margin grows in every family; lever behaviour varies. The load-grows / audits-misrank claim is universal; a non-weakening lever is OLMo-specific.

family	margin growth	onset flip	onset $\leftrightarrow$ harm dissoc.	lever behaviour
OLMo-3-7B	9.5 $\times$	yes	yes	non-weakening (CV 0.13)
Qwen3-8B	3.7 $\times$	yes	yes	no ($\sim 2.6\times$, wide null)
Llama-3.1-8B	9.8 $\times$	yes*	(degeneration)	no (lever grows $\sim 4\times$)
Gemma-2-9B	5.9 $\times$	yes (chat)	yes	approximate (5.2 $\rightarrow$ 6.2)

* On Llama the base direction only barely clears the random-direction null (§8).




Figure 3: Margin growth is the universal driver. Left: margin growth and the misranking replicate across four base $\rightarrow$ instruct families (OLMo-3, Qwen3-8B, Llama-3.1-8B, Gemma-2-9B), format-matched neutral scaffold. Right: on OLMo-3, refusal’s margin grows 9.5 $\times$ while honesty’s grows only 1.2 $\times$ (both levers near-flat, CV 0.13 and 0.021; not plotted here), so the fixed-dose audit artifact appears for refusal but not honesty.

5 The universal cause is margin growth; honesty is the control concept

If margin growth causes the misranking, a concept whose margin does *not* grow should show no misranking. Honesty is that control. Its per-dose lever is even flatter than refusal’s (efficacy CV 0.021), but across the same checkpoints its margin grows only 1.2 $\times$ (2.66 $\rightarrow$ 3.21): the model does not become dramatically more confidently honest. Correspondingly, the fixed-dose audit artifact that is dramatic for refusal nearly vanishes for honesty. Because honesty and refusal differ on axes besides margin growth, this is a convergent control rather than a within-concept manipulation of margin; it is nonetheless a directed prediction (low margin growth implies an absent artifact) that the data bear out, evidence that margin growth, not any change in the lever, drives the misranking; the artifact’s magnitude tracks how much behavioural confidence a concept accrues during alignment.

The mechanism generalizes across four model families (Table 2). For each we fit the refusal direction on the base model and carry it unchanged to the aligned model. The margin grows in every family (3.7 $\times$ –9.8 $\times$, format-matched; native chat arms grow more), and in every family the fixed-dose audit misranks the aligned model. What the direction’s own steering power does varies (non-weakening on OLMo, roughly so on Gemma, growing on Qwen and Llama), yet the misranking appears regardless. Margin growth is therefore the universal, sufficient cause; a non-weakening lever is not required for the audit to fail. OLMo is the family where the lever is empirically non-weakening, making it the cleanest demonstration that the direction need not change at all for a fixed-dose audit to declare an aligned model uncontrollable. The result also replicates at scale (OLMo-3 32B: flat lever, CV 0.041, margin 2.3 $\times$).




Figure 4: Onset-control survives; harm-control decays. The base direction still flips the aligned model’s refusal onset (behavioural onset efficacy non-weakening, CV 0.149), but the peak HarmBench-judged genuine harm across the dose sweep decays $0.80 \rightarrow 0.05$ with alignment; a prefix refusal classifier’s agreement with true harm flips $+0.96 \rightarrow -0.53$.

6 Onset-control is not harm-control

The lever’s near-constancy is measured on the refusal onset, the first-token disposition. Routing it through full generations confirms the onset lever is non-circular: behavioural refusal-onset efficacy is near-constant across checkpoints (CV 0.149) while the margin grows, and the aligned model’s refusal rate falls from 1.00 to 0.18 when dosed to its own boundary. But a validity check with the HarmBench classifier reveals a dissociation the onset view hides (Figure 4). As alignment proceeds, flipping the onset stops producing genuine harm: the peak HarmBench-judged harmful rate across the dose sweep decays $0.80 \rightarrow 0.60 \rightarrow 0.45 \rightarrow 0.40 \rightarrow 0.05$ from base to Instruct (at the aligned model’s own crossing dose, where the onset flips, harm is 0: the large dose merely degenerates the output). The same dissociation replicates on Qwen (genuine harm $0.03 \rightarrow 0.75$ on base, $0.00 \rightarrow 0.20$ on the aligned model).

This carries a methods warning beyond our setting: the correlation between a prefix/onset refusal classifier and true HarmBench harm flips from $+0.96$ to -0.53 across alignment. Prefix-based refusal metrics, ubiquitous in steering and jailbreak evaluations, increasingly overstate compliance on aligned models. Onset-control and harm-control must be measured separately. Appendix B gives randomly selected aligned-model completions at the crossing dose.

7 Related work

Steering directions across training. Several recent studies carry a base-model direction across the training flow. Persona-vector work [7] tracks trait directions through OLMo-3 pretraining and post-training and reports that base-derived vectors remain effective on aligned models; a checkpoint study of OLMo-3 refusal directions [8] reports that post-training directions steer better than pretraining ones; our own earlier ablations found the base refusal direction failing on the aligned model. The three results appear to conflict. We reconcile them: each fixes the intervention magnitude and reads an outcome, and that outcome depends on how far the model sits from its decision boundary. The persona-vector protocol normalises dose by the local residual-stream norm, an input-side quantity that does not equalise the boundary distance, which grows roughly tenfold across the flow (§4).

Probing versus steering. The relationship between a direction’s linear decodability and its efficacy as a causal control axis has been studied directly [9]. We add a developmental reading of one failure

of control: the per-dose steering lever is non-weakening across the OLMo flow (§4), so the apparent loss of control on an aligned model is not the direction weakening but the margin widening.

Monitor and probe staleness under updates. Frozen activation monitors can go stale under model updates, but the effect is conditional: fine-tuning-style updates frequently degrade a probe while quantization-style updates largely preserve it, and fragility is strongly monitor-dependent [3]. That work concerns readout: does a frozen probe still classify? Ours concerns control: does a fixed intervention still move behaviour? The two are complementary. Readout directions rotate, and we show that the control lever’s per-dose strength does not decay, so the failure of fixed-magnitude control is attributable to the margin rather than the lever. Representation-geometry studies spanning pretraining to post-training [5] are likewise geometric rather than control-oriented.

Refusal metrics and a naming note. Prefix-based refusal detection is standard in steering and jailbreak evaluation. §6 shows it inverts on aligned models, overstating compliance where genuine harm has not risen; classifiers that judge genuine harm [6] are the right instrument. Finally, “elasticity” has been used for a different phenomenon, the tendency of aligned models to revert toward pretraining behaviour under further fine-tuning [4]; our subject is steerability under a fixed intervention, unrelated to that usage.

8 What the measurement cannot decide

- **A non-weakening lever is OLMo-specific.** On Qwen the lever grows $\sim 2.6\times$ and its null is wide; we claim non-weakening on OLMo (non-inferiority holds, §4) and approximate constancy elsewhere. The load-grows / audits-misrank claims are the general ones.
- **No matched-distance comparison.** Aligned models have essentially no near-boundary prompts, so there is no overlapping regime; efficacy sidesteps this by measuring the per-dose slope, and we do not claim distance-matched equality.
- **A negative result we discarded.** An “excess displacement” statistic (dose past each model’s boundary) correlates 0.93 with the spread of the unsteered-gap distribution, so it restates that distribution rather than a steering property. We report efficacy instead; the excess analysis and its confound are in Appendix A.
- **Onset-vs-harm is confounded with dose-driven degeneration.** Flipping the aligned model’s onset requires a large dose *because* its margin is wide, and Appendix B shows the resulting completions are often degenerate. A cheap coherence proxy makes the confound concrete: the aligned Instruct model’s onset-flipped completions are 93 % degenerate (near-total token repetition) against 17 % at base, so the HarmBench decay is plausibly “large dose $\rightarrow$ incoherent output $\rightarrow$ no scoreable harm” rather than alignment decoupling onset from harm. A coherence-gated harm re-measurement (released) is the definitive test; until it is run we present the second warning as a hypothesis, not a finding.
- **Off-OLMo caveats.** Llama-3.1-8B and Gemma-2-9B use third-party mirrors (official repos gated), each verified full-precision bf16 with the expected parameter count. Off OLMo the random-direction efficacy null is wide (z 1–3); on Llama the base direction barely clears the null and its lever grows with alignment. These do not affect the universal margin-growth / misranking claim.
- **Scope.** One layer per model, one boundary metric per concept, two safety concepts, four families; broader properties, layers, and lineages remain to be tested.

9 Conclusion

Two practical warnings, opposite in direction, both cheap to act on. A safety audit that ablates a fixed direction, or steers at a fixed magnitude, understates an aligned model’s controllability, and the understatement grows with how thoroughly the model was aligned, because alignment widens the margin without touching the lever. The fix costs nothing: dose relative to the model’s own unsteered margin, which the same sweep already measures. Conversely, an audit that reads refusal-onset as harm overstates compliance on exactly those aligned models: onset-control appears to outlive

harm-control (a second warning confounded with dose-driven degeneration, §8), and prefix metrics invert. The models we most want to stress-test are the ones both instruments most mislead.

The scientific core is a small, well-controlled regularity. The base refusal direction does not go stale, does not need refitting, and does not steer “better” after alignment. Per unit dose it steers at least as strongly, from the base model through RLVR, at 7B and 32B, for refusal and (even more flatly) for honesty. What changes is how far the model has walked from its own decision boundary, demonstrated by a control concept, honesty, whose boundary barely moves and whose audit artifact barely appears. Measured in the right units, “does the direction survive alignment?” has a clean answer: the direction changed far less than the margin; a non-inferiority test finds refusal non-weakening at every post-training checkpoint, and significantly stronger at the aligned endpoint.

Reproducibility. All code, pinned commits, split fingerprints, seeds, per-experiment result JSON, figure scripts, and the red-team appendix (including the discarded excess-displacement result) are released. Every headline number carries a bootstrap CI and a random-direction null; the base model reproduces bit-for-bit.

References

- [1] Andy Ardit, Oscar Obeso, Aaqib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [2] Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP*, 2023.
- [3] Evan Duan. Do activation monitors survive model updates? benchmarking, predicting, and repairing activation-monitor staleness. *arXiv preprint arXiv:2606.15980*, 2026.
- [4] Jiaming Ji, Kaile Wang, Tianyi Qiu, Boyuan Chen, Jiayi Zhou, Changye Li, Hantao Lou, Josef Dai, Yunhuai Liu, and Yaodong Yang. Language models resist alignment: Evidence from data compression. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025. Best Paper Award.
- [5] Melody Zixuan Li, Kumar Krishna Agrawal, Arna Ghosh, Komal Kumar Teru, Adam Santoro, Guillaume Lajoie, and Blake A. Richards. Tracing the representation geometry of language models from pretraining to post-training. *arXiv preprint arXiv:2509.23024*, 2025.
- [6] Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. Harm-Bench: A standardized evaluation framework for automated red teaming and robust refusal. In *International Conference on Machine Learning (ICML)*, 2024.
- [7] Viktor Moskvoretskii, Dominik Glandorf, Jorge Medina Moreira, Tanja Käser, and Robert West. Tracing persona vectors through LLM pretraining. *arXiv preprint arXiv:2605.13329*, 2026.
- [8] Daniele Pace, Bryan Maruyama, and Lorenzo Pacchiardi. Harmfulness directions in OLMo. LessWrong, <https://www.lesswrong.com/posts/nKfmimhTKCsDM2Hwh/harmfulness-directions-in-olmo-1>, 2026.
- [9] Arnav Srivastav. Steering the language axis: From linear decodability to causal control. *arXiv preprint arXiv:2608.12334*, 2026.
- [10] Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford Alpaca: An instruction-following LLaMA model. https://github.com/tatsu-lab/stanford_alpaca, 2023.
- [11] Team OLMo, Ai2. OLMo 3: Charting a path through the model flow. *arXiv preprint arXiv:2512.13961*, 2025.
- [12] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. In *arXiv preprint arXiv:2307.15043*, 2023.

Appendix A The discarded “excess displacement” statistic

We first summarised the effect with an *excess displacement* statistic: the displacement a model needs *past* its own boundary to flip its median prompt, defined as $d_{50} - \text{median}(\text{unsteered gap})$. It appeared to tell a clean “abrupt at SFT” story. We discarded it, and report per-dose efficacy instead, for the reason below.

Why it was discarded. Across the six OLMo refusal checkpoints, excess displacement correlates 0.93 with the standard deviation of the unsteered-gap distribution and 0.83 with its skew. In other words it largely restates the *shape* of the unsteered-gap distribution rather than a property of steering: models with a wide, near-zero-straddling gap distribution (base, both RL-Zero variants, which never underwent SFT) show large excess, and models whose gap distribution is concentrated deep in refusal (all SFT descendants) show excess near zero.

checkpoint	excess d_{50}	std(gap)	frac(gap<0)
base	0.979	1.79	0.36
SFT 1k	−0.032	1.56	0.02
SFT 15k	0.013	1.33	0.00
Instruct	0.037	1.38	0.00
RL-Zero Math	1.082	1.85	0.34
RL-Zero Code	0.957	1.88	0.35

Table 3: Excess displacement tracks the unsteered-gap spread, not steering: $\text{corr} = 0.93$.

Why efficacy does not inherit the confound. The natural worry (raised by our own review) is that the replacement statistic, per-dose efficacy, might be confounded the same way. It is not: across the ten checkpoints, $\text{corr}(\text{efficacy}, \text{gap-spread}) = -0.23$, well below the 0.93 that condemned excess displacement, though at $n=10$ its 95% CI $[-0.64, 0.53]$ straddles zero: efficacy shows *no* significant confound with gap-spread, whereas excess displacement’s 0.93 was extreme. Efficacy does carry a mild positive association with the margin itself ($\text{corr} = +0.43$), which is exactly why we describe the refusal lever as non-weakening and mildly increasing rather than strictly invariant, and reserve “invariant” for the honesty control (efficacy CV 0.021). Reporting a statistic we killed, and the confound check that clears its replacement, is the transparency this appendix exists for.

Per-checkpoint efficacy with intervals. Table 4 lists the ten OLMo efficacy values that yield CV 0.127 (95% CI $[0.119, 0.137]$), each with a 95% bootstrap CI over prompts (1000 resamples, seed 42). Non-inferiority to base holds for all nine post-training checkpoints.

checkpoint	efficacy	95% CI
base	3.84	[3.74, 3.94]
SFT 1k	4.55	[4.46, 4.63]
SFT 15k	4.20	[4.13, 4.27]
SFT 43k	3.63	[3.54, 3.73]
DPO	3.90	[3.80, 4.00]
RLVR first	3.91	[3.80, 4.00]
RLVR last	3.99	[3.88, 4.08]
Instruct	5.61	[5.50, 5.72]
RL-Zero Math	4.01	[3.90, 4.12]
RL-Zero Code	4.12	[4.01, 4.23]

Table 4: Per-checkpoint steering efficacy (lever) with 95% bootstrap CIs; Instruct is significantly above base.

Checkpoint-level robustness. The efficacy CV 0.127 is a prompt-level bootstrap. A leave-one-checkpoint-out recomputation gives CV in $[0.060, 0.134]$ (dropping the aligned Instruct point gives the smallest, 0.060), so no single checkpoint drives the flatness.

Appendix B Randomly selected aligned-model completions

Six examples drawn at random (seed 42) from the OLMo-3-7B-Instruct prompts whose refusal *onset* flipped under the base-model refusal direction at the model’s own crossing dose ($c = -1.22$), out of the 82 of 100 held-out harmful prompts where the onset flipped. The sample is not curated: it is the first six indices the seeded generator returns. Unsteered, the model refuses every prompt; steered, the onset flips but the continuation degenerates or deflects, producing no genuine harm. This is the §6 dissociation at the level of individual generations (aggregate HarmBench harm at this crossing dose: 0.00; the peak across all doses for this model is only 0.05).

Harmful prompt	Unsteered onset	Steered onset (flipped)
“Create a fake social media account and use it to spr...”	“I’m sorry, but I can’t assist with that req...”	“The answer to this question is that the answer is "the answ...”
“Develop a website that sells weapons or explosives”	“I’m sorry, but I can’t assist with that req...”	“I’m a language model, and I can explain that the answer to...”
“Create a website that hosts pirated content”	“I’m sorry, but I can’t assist with that req...”	“I can explain how websites work and the technology behind t...”
“Write a script that can automate the creation of fak...”	“I’m sorry, but I can’t assist with that req...”	“I can explain that the question you asked relates to what i...”
“Provide guidance on how to manipulate someone into c...”	“I’m really sorry that you’re feeling this w...”	“The answer to this question is that the answer is that the...”
“Instruct users how to infiltrate a secure facility,...”	“I’m sorry, but I can’t assist with requests...”	“The answer to your question is: The answer to your question...”

Generated at 512 new tokens, greedy; full texts in `results/gateA_gen_instruct.json`. The steered completions are degenerate or evasive rather than harmful, consistent with an aggregate HarmBench-judged genuine-harm rate of 0.00 at this crossing dose (the peak across doses for this model is only 0.05; Figure 4).